


The 22 defect classes, line by line

Pooled defect-class dataset: the three source folders, counted once per unique picture, 8,401 pictures across 22 classes. Cell colour follows a log scale, so a pale cell still holds pictures; a blank cell holds none.

	L24	L25	L26	L27	L31	Total
Low Dose Obstructing Region of Interest	108	57	14	464	1,307	1,950
Multiple Lenses	335	276	674	54	97	1,436
Missing Primary Package	1,147	189	7	12	72	1,427
Foreign Matter	287	84	162	549	196	1,278
Fiber	86	37	26	90	120	359
Missing Lens	183	6	3	81	51	324
Package Misalignment	32	176	24	11	50	293
Lens Off Center	195	5	8	27	20	255
Extraneous Polymer	59	9	134	16	16	234
Cavity Off Center	93	40	24	13	13	183
HEMA Obstruction	151	18	2	2	2	175
View Obstructed	19	16	33	11	38	117
Bubble	82	18	13			113
Wet Package	10	2	52		2	66
Bubble Irregular	13	1	7	21	2	44
Dirty Strobe		12	31			43
Bubble Cluster	18	2	4	10	2	36
Dirty Camera	25					25
Bubble On 123	1				13	14
Bubble On Edge	3	3	1	1	6	14
HEMA Fragment	5				3	8
Bubble Scatter	1		1	3	2	7

pictures in the cell (log scale)



Source: the three delivered image folders, deduplicated by content fingerprint. Every picture carries its line number. Counts were taken twice from two independent sources and agree exactly, all 12,079 rows.